

Day 2 – Indexing, Monitoring, Backup & Atlas (8 Hours)

1. MongoDB Indexes (1 hr)

- Why indexes matter
 - Single-field indexes
 - Query planner introduction
 - Lab: Create and test index performance
-

2. MongoDB Indexes in Detail (1 hr 30 min)

- Compound indexes
 - Multikey indexes
 - Text indexes
 - TTL indexes
 - Partial & Sparse indexes
 - Index performance monitoring
 - Lab: Identify slow queries & fix them with indexes
-

3. MongoDB Logging Basics (1 hr)

- Log file structure
 - Slow query logs
 - Profiling levels
 - Lab: Enable profiler and analyze logs
-

4. Getting Started with MongoDB Atlas (1 hr)

- Create free tier cluster
 - Cluster deployment options
 - Network & security configuration
 - Lab: Deploy and connect to Atlas cluster
-

5. MongoDB Administrator Tools (1 hr)

- mongosh
 - mongodump, mongorestore
 - mongotop, mongostat
 - Compass & Atlas Metrics
 - Lab: Use admin tools on a dataset
-

6. Database Metrics & Monitoring (1 hr 30 min)

- Understanding MongoDB metrics
 - WiredTiger cache
 - IOPS, CPU, memory metrics
 - Monitoring tools: Atlas, Cloud Manager, Prometheus, Grafana
 - Lab: Live monitoring of workload
-

End of Day Assignment

- Create indexes and measure performance gain
 - Generate and analyze profiler logs
-

1. MongoDB Indexes (1 hr)

- Why indexes matter
 - Single-field indexes
 - Query planner introduction
 - Lab: Create and test index performance
-

Indexes are *the most critical performance feature* in MongoDB. Good indexing can reduce query time from **seconds** → **milliseconds**.

A. Why Indexes Matter

MongoDB stores documents in collections **without any built-in ordering**.

To search efficiently, indexes are required.

1. What is an Index?

An index is a data structure stored separately that allows MongoDB to quickly locate documents.

Think of it like the **index of a book** — instead of scanning every page, you jump directly to the needed section.

2. Without Index (Collection Scan – COLLSCAN)

MongoDB checks **every document** in the collection → slow.

```
> COLLSCAN  
> nScannedDocs = 50,000
```

3. With Index (Index Scan – IXSCAN)

MongoDB uses a B-tree structure → fast.

```
> IXSCAN  
> nScannedDocs = 50
```

4. Why CSC Needs Indexes

Use cases inside Corporation Service Company (CSC):

✓ Compliance dashboards

Query by jurisdiction, entityName, dueDate

✓ Domain monitoring

Search by domain, expiry, DNS type

✓ Trademark queries

Find by class, country, renewal date

✓ Officer lookup

Query by name, role, country

Indexes ensure:

- Faster search
 - Reduced CPU
 - Lower latency for APIs
 - Better customer experience
-

5. Indexes Reduce:

- CPU load
 - Response time
 - Memory pressure
 - Disk I/O
-

6. Indexes Increase:

- Storage usage
 - Write cost (because indexes must be updated on every insert/update)
-

B. Single-Field Indexes (20 minutes)

Basic syntax:

```
db.collection.createIndex({ field: 1 })      // Ascending
db.collection.createIndex({ field: -1 })        // Descending
```

1. Ascending Index Example

```
db.employees.createIndex({ salary: 1 });
```

Useful for:

- Range queries

- Sorting
-

2. Descending Index Example

```
db.entities.createIndex({ lastFiled: -1 });
```

Useful when retrieving:

- Latest filing records
 - Latest audit trails
-

3. Unique Index

Prevents duplicate values.

```
db.users.createIndex({ email: 1 }, { unique: true });
```

CSC Use Case:

- Unique corporate entity IDs
 - Unique trademark numbers
 - Unique domain names
-

4. Partial Index

Index only documents matching a filter.

```
db.logs.createIndex(
  { status: 1 },
  { partialFilterExpression: { status: "ERROR" } }
);
```

CSC Use Case:

- Index only compliance errors
 - Index only active legal entities
-

5. Index Options

Option	Meaning
unique	Prevent duplicates
sparse	Only index documents where the field exists
expireAfterSeconds	TTL index for auto-expiry

Option	Meaning
name	Custom index name

Example – TTL Index

Delete audit logs after 30 days:

```
db.auditLogs.createIndex({ ts: 1 }, { expireAfterSeconds: 2592000 });
```

C. Query Planner Introduction (10 minutes)

MongoDB must decide *how* to execute a query.

This is called the **Query Planner**.

1. explain() Function

Shows how MongoDB runs a query.

```
db.employees.find({ city: "Pune" }).explain("executionStats");
```

Important fields:

Field	Meaning
COLLSCAN	Slow, no index used
IXSCAN	Index used
nReturned	Number of returned documents
nScannedDocuments	How many scanned before match
executionTimeMillis	Query speed

2. Example – Without Index

```
db.customers.find({ city: "Delhi" }).explain("executionStats");
```

Output:

```
stage: "COLLSCAN"
nScannedDocuments: 10000
executionTimeMillis: 52
```

3. Example – With Index

```
db.customers.createIndex({ city: 1 });

db.customers.find({ city: "Delhi" }).explain("executionStats");
```

Output:

```
stage: "IXSCAN"
nScannedDocuments: 10
executionTimeMillis: 1
```

Interpretation

- The index drastically improves performance
 - Only a few documents scanned
 - Planner prefers indexed paths
-

D. LAB — Create & Test Index Performance (15 minutes)

Goal:

Observe difference between **COLLSCAN** and **IXSCAN**.

Step 1 — Create lab database & insert dataset

```
use csc_index_lab;

for (let i = 1; i <= 30000; i++) {
  db.entities.insertOne({
    entityId: i,
    entityName: "Entity " + i,
    jurisdiction: ["DEL", "NY", "TX", "AMS"][i % 4],
    lastFiled: new Date(2024, (i % 12), (i % 28) + 1)
  });
}
```

Step 2 — Query without index

```
db.entities.find({ jurisdiction: "DEL" }).explain("executionStats");
```

Expected:

```
COLLSCAN
nScannedDocuments: 30000
executionTimeMillis: ~20ms
```

Step 3 — Create Index

```
db.entities.createIndex({ jurisdiction: 1 });
```

Step 4 — Query again

```
db.entities.find({ jurisdiction: "DEL" }).explain("executionStats");
```

Expected:

```
IXSCAN
nScannedDocuments: ~7500
executionTimeMillis: ~2ms
```

Step 5 — Test sort performance

Without index:

```
db.entities.find().sort({ lastFiled: -1 }).explain("executionStats");
```

Create index:

```
db.entities.createIndex({ lastFiled: -1 });
```

Re-run sort:

```
db.entities.find().sort({ lastFiled: -1 }).explain("executionStats");
```

LAB Outcomes

Students will:

- Understand how indexes reduce scan time
 - Compare COLLSCAN vs IXSCAN
 - Learn how query planner chooses paths
 - Experience index performance effects in real-time
-

2. MongoDB Indexes in Detail (1 hr 30 min)

This module builds on basic indexing and covers:

- Compound Indexes
- Multikey Indexes
- Text Indexes
- TTL Indexes
- Partial & Sparse Indexes
- Index Monitoring & Query Performance

- Hands-on Index Optimization Lab
-

A. Compound Indexes (20 min)

A **compound index** includes **multiple fields** in a specific order.

✓ Syntax

```
db.collection.createIndex({ field1: 1, field2: -1 })
```

Key Rule: Index ORDER matters

Index: { a: 1, b: 1 }

Can support queries on:

- a
- a + b

Cannot support:

- b alone
-

★ CSC Use Case:

Querying legal entities:

```
db.entities.find({  
    jurisdiction: "DEL",  
    entityName: /^CSC/  
});
```

Optimal index:

```
db.entities.createIndex({ jurisdiction: 1, entityName: 1 });
```

Sort Optimization

If your query uses:

```
sort({ entityName: 1 })
```

AND the filter is:

```
{ jurisdiction: "DEL" }
```

Then compound index:

```
{ jurisdiction: 1, entityName: 1 }
```

removes in-memory sorting → faster.

! WRONG INDEX ORDER

Don't create this:

```
{ entityName: 1, jurisdiction: 1 }
```

Because the query filters by jurisdiction first → index won't help.

★ Compound Index Prefix Rule

Index { a: 1, b: 1, c: 1 }

Supports prefix queries:

- a
- a + b
- a + b + c

Does NOT support:

- b
- c
- b + c

B. Multikey Indexes (10 min)

Used for **array fields**.

✓ Automatically created when indexing an array

```
db.products.createIndex({ tags: 1 })
```

If document:

```
{ tags: ["electronics", "gaming", "laptop"] }
```

MongoDB indexes **each value**.

★ CSC Example

Index officers:

```
db.entities.createIndex({ "officers.role": 1 });
```

Supports:

```
db.entities.find({ "officers.role": "Director" })
```

! Limitations:

- Cannot create compound index with more than one array field
 - Impacts write performance due to multiple index entries
-

C. Text Indexes (15 min)

Used for **full-text search**.

✓ Create a Text Index

```
db.products.createIndex({ description: "text" });
```

✓ Search Text

```
db.products.find({ $text: { $search: "gaming laptop" } });
```

★ CSC Use Cases

- Trademark description search
 - Legal entity notes search
 - Domain WHOIS text search
-

✓ Score Results

Sort by relevance:

```
db.products.find(
```

```
{ $text: { $search: "gaming" } },
{ score: { $meta: "textScore" } }
).sort({ score: -1 });
```

! Limitations

- One text index per collection
 - Not suitable for large-scale search (use Atlas Search instead)
-

D. TTL Indexes (10 min)

Automatically deletes documents after a fixed time.

✓ Syntax

```
db.auditLogs.createIndex({ ts: 1 }, { expireAfterSeconds: 86400 });
```

★ CSC Use Case:

Delete compliance logs 90 days after creation.

```
db.compliance.createIndex(
  { createdAt: 1 },
  { expireAfterSeconds: 7776000 }
);
```

TTL is perfect for:

- Session expiry
 - Temp files
 - Cache
 - Audit logs
-

! Warnings:

- TTL deletion is not immediate (checked every 60 seconds)
 - TTL does NOT work on capped collections
 - Field must be a **Date** type
-

E. Partial & Sparse Indexes (15 min)

1. Sparse Index

Indexes only documents where field exists.

```
db.customers.createIndex({ email: 1 }, { sparse: true });
```

Use cases at CSC:

- Entities missing tax IDs
- Optional compliance fields

Sparse index reduces index size dramatically.

2. Partial Index

Indexes only documents matching a filter.

```
db.customers.createIndex(
  { premiumScore: 1 },
  { partialFilterExpression: { premiumCustomer: true } }
);
```

CSC Examples:

Index only:

- Active legal entities
- High-priority trademarks
- Expiring domains

Example:

```
db.domains.createIndex(
  { expiry: 1 },
  { partialFilterExpression: { status: "ACTIVE" } }
);
```

F. Index Performance Monitoring (10 min)

★ 1. List Indexes

```
db.collection.getIndexes()
```

★ 2. Validate Index Usage via explain()

```
db.entities.find({ jurisdiction: "TX" })
  .explain("executionStats")
```

Look for:

- IXSCAN
 - nReturned
 - nScannedDocuments
 - executionTimeMillis
-

★ 3. Check Index Size

```
db.collection.stats().indexSizes
```

★ 4. Drop Unused Indexes

```
db.collection.dropIndex("index_name")
```

★ 5. Query Profiler

Enable profiler:

```
db.setProfilingLevel(1);
```

Find slow queries:

```
db.system.profile.find().sort({ millis: -1 }).limit(5);
```

★ 6. Atlas Performance Tools (If using Atlas)

- Index suggestions
 - Query profiler
 - Slow query alerts
-

G. LAB — Identify Slow Queries & Fix Them with Indexes (20 min)



Goal:

Create a realistic scenario with slow queries → analyze → index → optimize.

Step 1 — Insert 50,000 Test Records

```
use csc_index_advanced;

for (let i = 1; i <= 50000; i++) {
  db.logs.insertOne({
    logId: i,
    user: ["admin", "ops", "api", "system"][i % 4],
    status: ["INFO", "WARN", "ERROR"][i % 3],
    ts: new Date(2024, (i % 12), (i % 28)),
    details: "System event " + i
  });
}
```

Step 2 — Run a Slow Query

```
db.logs.find({ user: "admin", status: "ERROR" })
  .sort({ ts: -1 })
  .limit(5)
  .explain("executionStats");
```

Expected:

```
COLLSCAN
nScannedDocuments: 50000
executionTimeMillis: HIGH
```

Step 3 — Create a Compound Index

Optimal index:

```
db.logs.createIndex({ user: 1, status: 1, ts: -1 });
```

Step 4 — Re-run Query

```
db.logs.find({ user: "admin", status: "ERROR" })
  .sort({ ts: -1 })
  .limit(5)
  .explain("executionStats");
```

Expected:

```
IXSCAN
```

```
nScannedDocuments: VERY LOW  
executionTimeMillis: LOW
```

Step 5 — Analyze Improvement

Students compare:

- Before vs after index
 - nScannedDocuments
 - executionTimeMillis
 - Query planner stage
-

LAB Outcomes:

Learners will learn:

- How to identify slow queries
 - How to choose correct index fields
 - How index order impacts performance
 - How to validate improvements with explain()
 - How to fix real-world performance issues
-

3. MongoDB Logging Basics (1 hr)

- Log file structure
- Slow query logs
- Profiling levels
- Lab: Enable profiler and analyze logs

Logging is a core part of MongoDB database administration.

Admins use logs to identify performance issues, security events, slow queries, replication events, and operational failures.

A. Log File Structure

MongoDB writes logs into **mongod.log**, typically located at:

- Linux: /var/log/mongodb/mongod.log
 - Windows: C:\Program Files\MongoDB\Server\log\mongod.log
 - Atlas: Integrated logging via UI
-

1. Log Format

A typical log line looks like:

```
2025-01-21T12:45:32.123+0530 I COMMAND [conn82] command test.users
command: find { filter: { city: "Bangalore" }, limit: 1 }
planSummary: COLLSCAN keysExamined:0 docsExamined:50000
executionTimeMillis: 40
```

Key components:

Field	Meaning
Timestamp	When the event occurred
Severity	I (info), W (warning), E (error), F (fatal)
Component	COMMAND, STORAGE, NETWORK, SHARDING
Connection ID	Session identifier
Event	What operation was executed

2. Common Log Types

- **Startup logs** → version, storage engine, config
 - **Connection logs** → clients connecting/disconnecting
 - **Query execution logs** → slow queries
 - **Replication logs** → oplog events
 - **Crash & shutdown logs**
-

3. CSC Administrator Use Cases

Logging helps CSC DBAs:

- Detect slow dashboard queries
- Identify long-running compliance data scans
- Trace unauthorized access attempts
- Debug Atlas function/apis

- Monitor replication lag in multi-regional setups
-

B. Slow Query Logs (10 minutes)

A “slow query” is any query that exceeds the threshold defined by:

```
operationProfiling:  
  slowOpThresholdMs: 100
```

Default: **100 milliseconds**

MongoDB automatically logs slow queries into mongod.log.

★ Example Slow Query Log

```
I COMMAND [conn89] query test.entities  
filter: { jurisdiction: "TX" }  
planSummary: COLLSCAN docsExamined: 50000  
keysExamined: 0  
executionTimeMillis: 120
```

★ What you learn from a slow query log

- Query filter
 - Execution plan summary (COLLSCAN or IXSCAN)
 - Number of documents scanned
 - Execution time
 - Index usage
-

★ How to Reduce Slow Queries

- Add appropriate index
 - Rewrite inefficient filters
 - Avoid regex prefix searches
 - Improve projection
 - Use compound indexes
-

C. Profiling Levels (15 minutes)

MongoDB profiler helps capture detailed information about database activity.

★ There are 3 profiler levels:

Level	Meaning
0	Off (default). No queries logged.
1	Log slow operations only.
2	Log all operations (high overhead).

★ Enable Profiling Level 1 (Recommended)

Logs operations slower than default slowms (100ms):

```
db.setProfilingLevel(1);
```

★ Customize slowms Threshold

```
db.setProfilingLevel(1, { slowms: 50 });
```

★ Enable Profiling Level 2 (Debug mode)

```
db.setProfilingLevel(2);
```

⚠ Warning: High overhead. Use only during debugging.

★ Disable Profiling

```
db.setProfilingLevel(0);
```

★ View Profiler Output

Profiler stores its logs in:

```
system.profile
```

Use:

```
db.system.profile.find().pretty();
```

⭐ Example Profiling Document

```
{
  "op": "query",
  "ns": "test.entities",
  "command": { "find": "entities", "filter": { "jurisdiction": "NY" } },
  "keysExamined": 0,
  "docsExamined": 32000,
  "millis": 42,
  "planSummary": "COLLSCAN"
}
```

D. LAB — Enable Profiler & Analyze Logs (20 minutes)

Goal: Face a slow query. Enable profiler. Identify issue. Fix using an index.

Step 1 — Insert Test Dataset

```
use logging_lab;

for (let i = 1; i <= 40000; i++) {
  db.logs.insertOne({
    eventId: i,
    user: ["system", "admin", "api"][i % 3],
    status: ["INFO", "WARN", "ERROR"][i % 3],
    ts: new Date(2024, (i % 12), (i % 28)),
    message: "Event log entry " + i
  });
}
```

Step 2 — Enable Profiling

```
db.setProfilingLevel(1, { slowms: 20 });
```

Step 3 — Trigger a Slow Query

```
db.logs.find({ user: "admin", status: "ERROR" })
```

```
.sort({ ts: -1 })
.limit(5);
```

Step 4 — Read Profile Logs

```
db.system.profile.find().sort({ millis: -1 }).limit(3).pretty();
```

Look for:

- COLLSCAN
 - High millis
 - Many docsExamined
-

Step 5 — Identify Missing Index

Query needs index on:

```
{ user, status, ts }
```

Step 6 — Create Compound Index

```
db.logs.createIndex({ user: 1, status: 1, ts: -1 });
```

Step 7 — Rerun Query & Check Profiling

```
db.logs.find({ user: "admin", status: "ERROR" })
    .sort({ ts: -1 })
    .limit(5)
    .explain("executionStats");
```

Expected:

```
IXSCAN
docsExamined small
executionTimeMillis very low
```

Step 8 — Validate Profiler Again

```
db.system.profile.find().sort({ ts: -1 }).limit(3).pretty();
```

Compare profile before vs after index.

LAB Outcomes

Students will learn to:

- Enable MongoDB profiler safely
 - Analyze slow queries
 - Interpret planSummary (COLLSCAN vs IXSCAN)
 - Identify missing indexes
 - Resolve performance issues
 - Validate improvements using profiler
-

Getting Started with MongoDB Atlas (1 hr)

- Create free tier cluster
 - Cluster deployment options
 - Network & security configuration
 - Lab: Deploy and connect to Atlas cluster
-

12. Getting Started with MongoDB Atlas (1 hr)

MongoDB Atlas is MongoDB's **fully managed cloud database service**, available on AWS, Azure, and GCP.

Admins use Atlas for **production workloads**, scaling, backup, global deployment, and monitoring.

A. Create Free Tier Cluster (15 min)

Atlas provides an **M0 Free Tier cluster** that is ideal for learning and prototyping.

Step-by-Step Instructions

Step 1 — Sign Up

Go to:

<https://www.mongodb.com/cloud/atlas>

Sign up using:

- Google account, or
 - Email + password
-

Step 2 — Create New Project

Atlas organizes clusters inside **projects**.

Steps:

1. Click **New Project**
2. Enter name:

CSC-Mongo-Lab

3. Add members (optional)
-

Step 3 — Create Cluster

Select:

- **M0 Free Tier**
- Choose cloud provider: **AWS / Azure / GCP**
- Choose region nearest to India:

AWS ap-south-1 (Mumbai) → recommended

Step 4 — Choose Cluster Name

Default: Cluster0

Recommended name:

csc-training-db

Step 5 — Create Database User

Create a user for client connections:

Username:

`cscAdmin`

Password:

Choose strong password

Role:

readWriteAnyDatabase (for training)

Step 6 — Add IP Access List

To allow connections, add:

0.0.0.0/0

(Allows access from any IP — OK for lab, not recommended for production.)

Step 7 — Wait for Cluster Deployment

Takes 3–5 minutes.

B. Cluster Deployment Options (10 min)

Atlas supports multiple deployment sizes and architectures.

1. M0/M2/M5 (Shared Tier)

- Free or low-cost
 - Shared instances
 - No dedicated RAM/CPU
 - Good for learning & prototypes
-

2. Dedicated Clusters (M10 and above)

- Full control of compute
 - Production-ready
 - Supports sharding & backup
-

3. Global & Multi-region Deployment

Used by enterprises like CSC for:

- Low latency across teams
- Disaster recovery
- Data residency compliance

Modes:

- Multi-region write
 - Multi-region read
 - Custom regional failover
-

4. Advanced Features

- Auto-scaling compute & storage
 - Serverless instances
 - Atlas Search
 - Atlas Device Sync
 - Online archive
-

C. Network & Security Configuration (15 min)

Security is a top priority in Atlas.

1. IP Access List (Firewall)

Atlas allows connections *only* from whitelisted IP addresses.

Add:

Add My Current IP

or for lab:

0.0.0.0/0

2. Database User Authentication

RBAC-based permissions:

- read
- readWrite
- dbAdmin
- clusterAdmin

Creation UI:

Database Access → Add New Database User

3. TLS/SSL Encryption

Atlas **forces TLS connections**, ensuring secure communication.

Connection string includes:

`tls=true`

4. Network Peering (for enterprises)

Connect Atlas to:

- AWS VPC
- Azure VNet
- GCP VPC

Used by CSC for:

- Secure private connection
 - Isolation from internet
 - Compliance requirements
-

5. Encryption at Rest (Automated)

Atlas automatically encrypts storage using:

- AWS KMS
- Azure Key Vault

- GCP KMS
-

D. LAB — Deploy & Connect to Atlas Cluster (20 min)



Goal:

Learners should be able to create a cluster, configure security, and connect using:

- **MongoDB Compass**
- **mongosh**
- **Application driver** (Node.js/Python example)

LAB PART 1: Create Free Tier Cluster

Follow steps from Section A:

1. Create cluster
 2. Add IP whitelist
 3. Create DB user
-

LAB PART 2: Connect Using MongoDB Compass

Step 1 — Open Compass

Download from:

<https://www.mongodb.com/try/download/compass>

Step 2 — Get Connection String

In Atlas UI:

Connect → Compass → Copy connection string

Example:

`mongodb+srv://cscAdmin:<password>@csc-training-db.abcde.mongodb.net/test`

Replace <password> with your actual password.

Step 3 — Connect

Paste connection string into Compass and click **Connect**.

Students should be able to:

- View databases
- Add sample data
- Run queries

LAB PART 3: Connect Using mongosh

Step 1 — Install mongosh (if needed)

<https://www.mongodb.com/try/download/shell>

Step 2 — Connect to Atlas

```
mongosh "mongodb+srv://cscAdmin:<password>@csc-training-db.abcde.mongodb.net"
```

Step 3 — Create database & collection

```
use atlas_lab;  
db.products.insertOne({ name: "Laptop", price: 75000 });
```

LAB PART 4: Connect Using Application Driver

Node.js Example

```
const { MongoClient } = require("mongodb");  
  
const uri = "mongodb+srv://cscAdmin:<password>@csc-training-db.abcde.mongodb.net/";  
  
const client = new MongoClient(uri);  
  
async function run() {  
  await client.connect();  
  const db = client.db("atlas_lab");  
  console.log(await db.listCollections().toArray());  
}  
  
run();
```

LAB Outcomes

Students will be able to:

- Deploy & configure Atlas cluster
 - Understand free vs dedicated vs global clusters
 - Configure IP access, DB users, TLS
 - Connect via Compass, mongosh, and drivers
 - Start using Atlas as their primary MongoDB environment
-

MongoDB Administrator Tools

- mongosh
- mongodump, mongorestore
- mongotop, mongostat
- Compass & Atlas Metrics
- Lab: Use admin tools on a dataset

MongoDB administrators rely on a combination of **shell tools**, **diagnostic utilities**, and **UI tools** to manage, monitor, and troubleshoot clusters.

This module covers the core admin tools used in **on-prem, self-managed, and Atlas environments**.

A. mongosh – The MongoDB Shell (10 min)

mongosh is the modern JavaScript-based MongoDB shell.

Common Admin Commands

1. Connect to MongoDB

```
mongosh "mongodb://localhost:27017"
```

2. Show databases

```
show dbs
```

3. Show collections

```
use cscdb
show collections
```

4. Add admin user

```
db.createUser({
  user: "cscAdmin",
  pwd: "StrongPass123",
  roles: ["root"]
});
```

5. Server status (monitoring snapshot)

```
db.serverStatus()
```

★ CSC Use Cases

- Verify authentication status
 - Check replication state
 - Inspect server health before deployments
 - Run automated scripts for compliance data loads
-

B. mongodump & mongorestore (15 min)

These tools are essential for **backup**, **migration**, and **offline restore**.

1. mongodump (Backup)

Creates a BSON dump of the database.

Dump entire database

```
mongodump --db cscdb --out /backup/cscdb-backup
```

Dump a collection

```
mongodump --db cscdb --collection users --out /backup/users-backup
```

Dump Atlas cluster

```
mongodump --uri "mongodb+srv://<user>:<pw>@cluster.mongodb.net"
```

2. mongorestore (Restore)

Restores BSON backup to a MongoDB instance.

Restore a database

```
mongorestore --db cscdb /backup/cscdb-backup/cscdb
```

Restore a collection

```
mongorestore --collection users /backup/users-backup/cscdb/users.bson
```

★ CSC Use Cases:

- Migrating client legal entity data
 - Creating training/staging environments
 - Restoring archived compliance datasets
 - Backing up before major schema changes
-

C. mongotop & mongostat (15 min)

These tools help admins monitor **real-time activity**.

1. mongotop

Shows how much time the database spends reading/writing per collection.

Run:

```
mongotop 3
```

Output refreshed every **3 seconds**.

Interpretation

Column	Meaning
ns	Namespace (db.collection)
command	Query time
update	Update time
read/write	Time spent in disk ops

★ Use Cases at CSC:

- Detect hot collections (e.g., compliance events)
 - Identify collections causing disk pressure
 - Debug indexing spikes
-

★ 2. mongostat

Shows key server metrics.

Run:

```
mongostat 2
```

Columns include:

- inserts / updates / deletes
- qr | qw (queued reads/writes)
- ar | aw (active reads/writes)
- netIn / netOut
- connections
- % dirty / % used (WiredTiger cache)

★ What DBAs look for

Symptom	Meaning
High queued writes (qw)	Disk bottleneck
High connections	App connection leak
Cache dirty > 40%	Slow checkpoints
inserts spikes	Batch jobs or migrations

D. Compass & Atlas Metrics (10 min)

MongoDB Compass and Atlas UI provide graphical admin insights.

MongoDB Compass

Useful for:

- Browsing databases
 - Running queries
 - Schema visualization
 - Index analysis
 - Real-time performance charts
-

Atlas Metrics Dashboard

Atlas provides:

- CPU usage
 - Memory usage
 - IOPS
 - Connections
 - Replication lag
 - Slow query panel
 - Performance advisor (index recommendations)
-

★ CSC-Relevant Scenarios:

- Monitor ingestion workloads
 - Validate health post-deployment
 - Track performance before/after schema or index changes
 - Investigate spikes in client-facing apps
-

E. LAB – Use Admin Tools on a Dataset (10 min)

Lab Goal:

Learner should be able to:

- Import dataset
 - Inspect with mongostat/mongotop
 - Perform backup/restore
 - View metrics in Compass or Atlas
-

Step 1 — Load Sample Dataset

Create a test collection:

```
use admintools;
for (let i = 1; i <= 20000; i++) {
  db.events.insertOne({
    eventId: i,
    type: ["INFO", "WARN", "ERROR"][i % 3],
    user: ["admin", "system", "api"][i % 3],
    ts: new Date(),
    msg: "Event log " + i
  });
}
```

Step 2 — Run mongostat

```
mongostat 2
```

Observe:

- inserts increasing
 - connections
 - opcounters
-

Step 3 — Run mongotop

```
mongotop 3
```

Expected:

admintools.events shows high read/write activity.

Step 4 — Run mongodump / mongorestore

Backup:

```
mongodump --db admintools --out /tmp/backup
```

Drop collection:

```
db.events.drop();
```

Restore:

```
mongorestore --db admintools /tmp/backup/admintoools
```

Check:

```
db.events.countDocuments()
```

Step 5 — Open in Compass

- Connect
- Validate restored data
- Check schema tab

- Inspect index usage
-

Step 6 — Atlas Metrics (If using Atlas)

- Navigate to Metrics
 - View CPU spikes
 - Check connections
 - Open Profiler
-

LAB Outcomes:

Learners will understand:

- Backup & restore
 - Real-time server monitoring
 - Performance diagnostics
 - Using both command-line + GUI tools
-

Following Topics should be included with Lab Example

Database Metrics & Monitoring

- Understanding MongoDB metrics
- WiredTiger cache
- IOPS, CPU, memory metrics
- Monitoring tools: Atlas, Cloud Manager, Prometheus, Grafana
- Lab: Live monitoring of workload

Below is your **full instructor-ready Module 14: Database Metrics & Monitoring (1 hr 30 min)** for the MongoDB Administrator Course.

Database Metrics & Monitoring (1 hr 30 min)

Database monitoring is the backbone of MongoDB administration.

As workloads scale, admins must track:

- Latency

- Throughput
 - Cache health
 - Disk I/O
 - Replication lag
 - Query efficiency
-

A. Understanding MongoDB Metrics (20 min)

MongoDB exposes multiple categories of metrics through **serverStatus**, **collStats**, and monitoring tools.

1. Key Database Metrics

Latency Metrics

- Query latency (ms)
- Write latency
- Lock wait time

Throughput Metrics

- ops/sec → inserts, queries, updates
- network in/out
- queued operations

Storage Engine Metrics

- WiredTiger cache usage
- Checkpoint activity
- Disk I/O latency
- Dirty vs clean pages

Replication Metrics

- Oplog application time
- Replication lag
- Elections

2. How to check metrics manually

Query serverStatus:

```
db.serverStatus()
```

Collection-level metrics:

```
db.collection.stats()
```

WiredTiger metrics:

```
db.serverStatus().wiredTiger
```

3. CSC Enterprise Use Cases

Monitoring helps CSC administrators:

- Identify slow legal search queries
 - Detect spikes during client onboarding
 - Track compliance event ingestion load
 - Diagnose Atlas cluster cost drivers
 - Prevent replication lag across global offices
-

B. WiredTiger Cache (20 min)

WiredTiger is the default storage engine in modern MongoDB versions.

★ 1. What is WiredTiger Cache?

It is the in-memory cache used for:

- Document reads
- Index reads
- Write buffering
- Page management

Default size:

50% of RAM for standalone

25% of RAM for sharded clusters

★ 2. Key cache metrics

Check cache metrics:

```
db.serverStatus().wiredTiger.cache
```

Important fields:

Metric	Meaning
bytes currently in cache	Amount of data in RAM
maximum bytes configured	Cache size limit
pages read into cache	Read load
pages written from cache	Write load
% cache dirty	Dirty pages waiting flush
checkpoint blocked page eviction	Checkpoint causing stalls

★ 3. Indicators of a healthy cache

- ✓ Cache usage < 80%
- ✓ Dirty pages < 20–25%
- ✓ Page read/write ratio stable
- ✓ No checkpoint stalls

★ 4. Symptoms of cache issues

Symptom	Cause	Fix
Dirty pages > 40%	Heavy writes	Add index, optimize writes, scaling
Many page evictions	Cache thrashing	Increase instance size
Read stalls	Poor indexes	Create appropriate indexes
Checkpoint slow	Slow disk	Use SSD, adjust cache

★ 5. CSC Example

During quarterly client data ingestion, CSC team notices:

- Cache dirty pages 45%
- Slow dashboard queries

Action:

Add index on { clientId, createdAt } → queries become 8× faster.

C. IOPS, CPU, Memory Metrics (15 min)

★ 1. CPU Metrics

- % CPU usage
 - CPU spikes → caused by sorting, aggregations, or full scans
 - High “user” CPU → queries
 - High “system” CPU → OS-level overhead
-

★ 2. Memory Metrics

- WiredTiger cache usage
 - Resident memory
 - Queued read/write operations
 - Page faults → insufficient memory
-

★ 3. Disk Metrics (IOPS)

IOPS = Operations Per Second

Measure:

- read IOPS
- write IOPS
- fsync IOPS during checkpoints

High IOPS usage indicates:

- Index building
 - Large aggregations
 - Update-heavy workloads
-

★ 4. How to identify IOPS problems

Look for:

`mongostat`

Indicators:

- High “flushes”
 - High “dirty” percentage
 - High queue length
-

D. Monitoring Tools (Atlas, Cloud Manager, Prometheus, Grafana) (20 min)

★ 1. MongoDB Atlas Monitoring

Atlas provides built-in dashboards:

Key charts in Atlas:

- CPU
- Memory
- IOPS
- Connections
- Opcounters
- Replication lag
- Slow queries
- Index suggestions

Atlas Profiler

Shows:

- Slowest queries
 - Query shapes
 - Execution plans
 - Index suggestions
-

★ 2. Cloud Manager (Self-managed deployments)

Cloud Manager provides:

- Automation
- Backups
- Monitoring
- Alerts

Used for:

- On-prem MongoDB clusters
 - Hybrid environments
-

★ 3. Prometheus + Grafana

These tools are used by DevOps/SRE teams.

Prometheus:

- Pulls metrics from MongoDB Exporter
- Stores time series metrics

Grafana:

- Visualizes dashboard panels:
 - CPU
 - Memory
 - Op counters
 - IOPS
 - Replication lag
 - Cache usage

★ Recommended Grafana Dashboards

- MongoDB Overview
 - WiredTiger Cache
 - Replication Performance
 - Cluster Health
-

E. LAB — Live Monitoring of Workload (15 min)



Goal:

Generate load → monitor metrics in real time → identify anomalies.

LAB PART 1 — Insert workload generator

Create 50k documents:

```
use monitorlab;

for (let i = 1; i <= 50000; i++) {
  db.logs.insertOne({
    logId: i,
    type: ["INFO", "WARN", "ERROR"][i % 3],
    user: ["system", "client", "admin"][i % 3],
    ts: new Date(),
    details: "Log event " + i
  });
}
```

LAB PART 2 — Monitor via mongostat

In terminal:

```
mongostat 2
```

Observe:

- inserts
 - queries
 - updates
 - connections
 - queue depth
-

LAB PART 3 — Monitor via mongotop

```
mongotop 3
```

Observe:

monitorlab.logs gaining heavy read/write time.

LAB PART 4 — Observe Cache Metrics

```
db.serverStatus().wiredTiger.cache
```

Look for:

- bytes in cache
 - dirty bytes
 - eviction rates
-

LAB PART 5 — Atlas Metrics (if using Atlas)

Open Metrics dashboard and observe:

- CPU %
 - IOPS
 - Network usage
 - Slow query panel
-

LAB PART 6 — Identify a Slow Query

Run:

```
db.logs.find({ type: "ERROR" }).sort({ ts: -1 }).limit(5);
```

View profiler:

```
db.system.profile.find().sort({ millis: -1 }).limit(1).pretty();
```

LAB PART 7 — Fix with Index

Create:

```
db.logs.createIndex({ type: 1, ts: -1 });
```

Re-run workload and compare metrics in:

- mongostat

- mongotop
 - serverStatus
 - Atlas
-

LAB Outcomes:

Students will:

- Understand how metrics react to workload
 - Interpret cache, CPU, IOPS indicators
 - Detect slow queries in real time
 - Use multiple monitoring tools together
 - Optimize using indexes
-

End of Day Assignment (Indexes + Profiler & Logs)

Objective

By the end of this assignment, the learner should be able to:

- Identify slow queries and **speed them up using indexes**
 - Use explain() to **measure performance gain**
 - Configure the **profiler**, generate logs, and analyze **slow queries**
-

PART A – Create Indexes & Measure Performance Gain (45–60 min)

1 Dataset Setup

Open mongosh and run:

```
use day2_perf_lab;

for (let i = 1; i <= 100000; i++) {
  db.orders.insertOne({
    orderId: i,
    customerId: "CUST" + (i % 5000),
    status: ["NEW", "SHIPPED", "DELIVERED", "CANCELLED"] [i % 4],
    amount: Math.floor(Math.random() * 9000) + 1000,
    country: ["IN", "US", "UK", "DE", "SG"] [i % 5],
    createdAt: new Date(2024, (i % 12), (i % 28) + 1)
});
```

```
}
```

2 Task 1 – Baseline Query (Without Index)

Run this query and capture the execution stats:

```
db.orders.find(
  { country: "IN", status: "DELIVERED" }
).sort({ createdAt: -1 }).limit(10).explain("executionStats");
```

Record in your notes:

- executionTimeMillis
 - executionStats.totalDocsExamined
 - executionStats.totalKeysExamined
 - executionStages.stage (COLLSCAN or IXSCAN?)
-

3 Task 2 – Design & Create Index

Based on the query pattern, design a **compound index**.

Suggested:

```
db.orders.createIndex(
  { country: 1, status: 1, createdAt: -1 }
);
```

4 Task 3 – Rerun Query (With Index)

Run the **same query** again with explain:

```
db.orders.find(
  { country: "IN", status: "DELIVERED" }
).sort({ createdAt: -1 }).limit(10).explain("executionStats");
```

Now record again:

- executionTimeMillis
- totalDocsExamined
- totalKeysExamined
- executionStages.stage

👉 **Compare Before vs After:**

Write a short note:

1. How much did executionTimeMillis change?

-
2. Did totalDocsExamined go down significantly?
 3. Did the plan change from COLLSCAN to IXSCAN?
 4. Is this index suitable for this query? Why?
-

Task 4 – Extra: Test Wrong Index Order (Optional)

Create another index:

```
db.orders.createIndex({ status: 1, country: 1, createdAt: -1 });
```

Now run the **same query** again with explain().

- Does MongoDB use the “better” index or the “wrong-order” index?
 - Check the winningPlan in explain() and identify **which index** was used.
-

✳ PART B – Generate & Analyze Profiler Logs (30–45 min)

Task 5 – Enable Profiler

Set the profiler to log operations slower than **20 ms**:

```
use day2_perf_lab;
db.setProfilingLevel(1, { slowms: 20 });
```

Verify:

```
db.getProfilingStatus();
```

Task 6 – Trigger Slow Queries

Run **two different queries**:

Query A – Less selective filter

```
db.orders.find({ status: "DELIVERED" }).sort({ createdAt: -1 }).limit(50);
```

Query B – More selective filter (indexed)

```
db.orders.find(
  { country: "IN", status: "DELIVERED" }
).sort({ createdAt: -1 }).limit(20);
```

Run each 2–3 times.

8 Task 7 – Read Profiler Output

Check the slowest operations:

```
db.system.profile.find().sort({ millis: -1 }).limit(5).pretty();
```

For at least **two entries**, note down:

- ns (namespace)
- command (filter + sort used)
- millis
- planSummary (COLLSCAN / IXSCAN)
- docsExamined / keysExamined (if present)

Answer:

1. Which query (A or B) appears slower in profiler?
 2. Does planSummary show COLLSCAN anywhere?
 3. Are any queries NOT using the index you created? Why might that be?
-

9 Task 8 – Fix a Slow Query Using Profiler Insight

1. From system.profile, pick **one slow query** that is not using an index optimally.
2. Design a new index for that query's filter/sort.
3. Create the index using createIndex().
4. Rerun the query & check profiler again.

Write a short summary:

- What was the original query?
- What index did you create?
- How did millis and docsExamined change?

Reflection Questions (Short Written Answers)

1. Why can a **compound index** be better than multiple single-field indexes?
 2. What is the difference between **COLLSCAN** and **IXSCAN** in explain()?
 3. Why should you be careful with setting profiler level to 2 in production?
 4. How can **profiler + indexes** together form a performance tuning workflow?
-