

MongoDB Administrator Lab

Day 1 – MongoDB Foundations & Core Administration

1. Introduction to MongoDB Ecosystem (1 hr)

- What is MongoDB?
 - Document Database vs Relational Databases
 - BSON & JSON structure
 - Collections, Databases, and Clusters
 - Where MongoDB fits in modern applications
-

2. MongoDB and the Document Model (1 hr)

- Embedded vs Referenced Documents
 - Flexible Schema
 - Sample document structures
 - Lab: Explore sample documents using MongoDB Shell
-

3. MongoDB Data Modelling Introduction (1 hr)

- Data modelling patterns
 - One-to-One, One-to-Many, Many-to-Many
 - Polymorphic patterns
 - Lab: Model a simple E-commerce schema
-

4. The MongoDB Shell (mongosh) (1 hr)

- Basic shell operations
 - Connecting to standalone, replica set
 - Shell queries and scripting basics
 - Lab: Connect and explore collections
-

5. Connecting to a MongoDB Database (30 min)

- Connection string concepts
- Users, roles, authentication
- Connect using mongosh, Compass, and drivers

6. MongoDB CRUD Operations – Insert & Find (1 hr)

- `insertOne()`, `insertMany()`
 - `find()`, `findOne()`
 - Query filters, projections
 - Lab: Perform CRUD queries on a sample dataset
-

7. MongoDB CRUD Operations – Replace & Delete (30 min)

- `replaceOne()`
 - `deleteOne()`, `deleteMany()`
 - Lab: Data cleanup operations
-

8. Modifying Query Results (1 hr)

- `sort()`, `skip()`, `limit()`
 - Aggregation preview
 - Query optimization basics
 - Lab: Apply query modifiers for performance
-

End of Day Assignment

- CRUD tasks on sample collections
- Short quiz (MCQ + practical)

LAB EXAMPLE – CSC Context (15 minutes hands-on)

Goal: Insert and query CSC Legal Entity data in MongoDB.

Step 1 – Create Database

```
use csc_legal
```

Step 2 – Create Collection

```
db.entities.insertOne({  
  entityName: "CSC Global Holdings LLC",  
  jurisdiction: "Delaware",  
  status: "Active",  
  officers: [  
    { name: "John Doe", role: "Director" },  
    { name: "Sarah Lee", role: "Secretary" }  
,  
  filings: {  
    annualReport: true,  
    lastFiled: ISODate("2024-09-15")  
  },  
  brandPortfolio: ["trademark", "domain"]  
});
```

Step 3 – Query

```
db.entities.find({ jurisdiction: "Delaware" }).pretty();
```

Step 4 – Filter by Officer Role

```
db.entities.find({ "officers.role": "Director" });
```

Step 5 – Projection Example

```
db.entities.find(  
  {},  
  { entityName: 1, jurisdiction: 1, _id: 0 }  
)
```

★ LAB Learning Outcomes

CSC administrators will learn:

- How MongoDB stores corporate/legal JSON structures
- How to insert and retrieve entity metadata

- How nested and array fields reflect real-world domain structures
-

2. MongoDB and the Document Model (1 hr)

- Embedded vs Referenced Documents
 - Flexible Schema
 - Sample document structures
 - Lab: Explore sample documents using MongoDB Shell
-

C. Sample Document Structures (10 minutes)

Below are real-world structures relevant to CSC's legal & compliance domain.

1. Corporate Entity Document

```
{  
  "entityId": "CSC-DEL-2025-001",  
  "entityName": "CSC Corporate Solutions Pvt Ltd",  
  "jurisdiction": "Delhi",  
  "officers": [  
    { "name": "Vikram Rao", "role": "Director", "since":  
      "2022" }  
  ],  
  "filings": {  
    "annualReport": true,  
    "lastFiled": "2024-07-15"  
  }  
}
```

2. Trademark Portfolio Document

```
{  
  "trademark": "CSC GLOBAL",  
  "classes": [9, 35, 42],  
  "countries": ["US", "UK", "IN"],  
  "status": "Registered",  
  "renewalDue": "2026-11-20"  
}
```

3. Digital Brand Asset Document

```
{  
  "domain": "cscglobal.com",  
  "dns": {  
    "A": "104.16.44.1",  
    "CNAME": "cscglobal.com.edge.net"  
  },  
  "sslExpiry": "2025-08-01"  
}
```

D. LAB: Explore Sample Documents using MongoDB Shell (15 minutes)

Goal: Insert and query CSC business documents using embedding and referencing.

Step 1 – Start the shell

```
mongosh
```

Step 2 – Create Database

```
use csc_training
```

Step 3 – Insert Embedded Document

```
db.entities.insertOne({  
  entityName: "CSC Global Holdings LLC",  
  jurisdiction: "Delaware",  
  officers: [  
    { name: "John Doe", role: "Director" },  
    { name: "Anna Smith", role: "Secretary" }  
  ]  
});
```

Step 4 – Insert Referenced Document

```
const log1 = db.auditLogs.insertOne({  
  action: "Created entity",  
  user: "admin",  
  ts: new Date()  
}).insertedId;  
  
db.entities.insertOne({  
  entityName: "CSC Digital Services India Pvt Ltd",  
  auditLogs: [log1]  
});
```

Step 5 – Run Queries

1. Find all entities

```
db.entities.find().pretty();
```

2. Find based on embedded officer field

```
db.entities.find({ "officers.role": "Director" });
```

3. Lookup referenced audit logs (with aggregation)

```
db.entities.aggregate([
  { $match: { entityName: "CSC Digital Services India Pvt Ltd" } },
  { $lookup: {
    from: "auditLogs",
    localField: "auditLogs",
    foreignField: "_id",
    as: "auditTrail"
  }
]) .pretty();
```

4. Projection Example

```
db.entities.find({}, { entityName: 1, jurisdiction: 1, _id: 0 });
```

Lab Learning Outcomes

Participants will understand:

- How MongoDB stores nested and related data
 - Difference between embedding and referencing
 - Querying embedded fields
 - Using aggregation \$lookup for joins
 - Real CSC-style structured documents
-

3. MongoDB Data Modelling Introduction (1 hr)

- Data modelling patterns
- One-to-One, One-to-Many, Many-to-Many
- Polymorphic patterns
- Lab: Model a simple E-commerce schema

D. LAB: Model a Simple E-Commerce Schema (15 minutes)

Students will design and insert **Products**, **Customers**, **Orders**, and **Order Items**.

Step 1 – Switch to Training DB

```
use ecommerce_demo
```

Step 2 – Insert Product Documents (Polymorphic Products)

```
db.products.insertMany([
  {
    _id: "P1001",
    type: "electronics",
    name: "Laptop",
    brand: "HP",
    price: 75000,
    warranty: 24
  },
  {
    _id: "P2001",
    type: "clothing",
    name: "T-Shirt",
    brand: "Nike",
    size: "M",
    price: 1200
  }
]);
```

Step 3 – Insert Customer

```
db.customers.insertOne({
  _id: "C001",
  name: "Amit Sharma",
  email: "amit@example.com",
  address: { city: "Pune", pincode: 411001 }
```

```
});
```

Step 4 – Insert Order (Embedded Items)

```
db.orders.insertOne({  
  _id: "O5001",  
  customerId: "C001",  
  orderDate: new Date(),  
  items: [  
    { productId: "P1001", qty: 1, price: 75000 },  
    { productId: "P2001", qty: 2, price: 1200 }  
  ],  
  totalAmount: 77400  
});
```

Step 5 – Query Examples

✓ Find Orders for a Customer

```
db.orders.find({ customerId: "C001" }).pretty();
```

✓ Find All Electronics Products

```
db.products.find({ type: "electronics" });
```

✓ Access embedded fields

```
db.orders.find({ "items.productId": "P2001" });
```

⌚ Lab Learning Outcomes

You will learn to:

- Build product models with polymorphism
 - Model 1-to-many using embedded items
 - Use referencing via productId
 - Query embedded arrays
 - Create e-commerce design patterns
-

MODULE 4: The MongoDB Shell (mongosh)

The MongoDB Shell (**mongosh**) is the primary interactive tool for DBAs and developers to connect, query, administer, and inspect MongoDB deployments.

A. Basic Shell Operations (20 minutes)

What is mongosh?

A JavaScript-powered shell used to:

- Connect to MongoDB deployments
 - Execute CRUD operations
 - Manage users, roles, indexes
 - Run administrative commands
 - Inspect performance and cluster metadata
-

1. Start the Shell

To start the shell (assuming default local installation):

```
mongosh
```

2. Show Databases

```
show dbs
```

Displays all databases with size > 0.

3. Create or Switch Database

```
use mydb
```

MongoDB creates the DB only after inserting the first document.

4. Show Collections

```
show collections
```

5. Insert a Document

```
db.users.insertOne({ name: "Amit", city: "Pune" });
```

6. Find Documents

```
db.users.find();
```

Pretty format:

```
db.users.find().pretty();
```

7. Update and Delete

```
db.users.updateOne({ name: "Amit" }, { $set: { city: "Bangalore" } });
db.users.deleteOne({ name: "Amit" });
```

8. Check Shell Help

```
help
db.help()
```

B. Connecting to Standalone & Replica Set (15 minutes)

1. Connect to a Standalone MongoDB Server

```
mongosh "mongodb://localhost:27017"
```

If username/password exists:

```
mongosh "mongodb://adminUser:P@ssw0rd@localhost:27017/admin"
```

2. Connect to a Replica Set

Replica set connection format:

```
mongodb://host1:27017,host2:27018,host3:27019/?replicaSet=rs0
```

Example:

```
mongosh
"mongodb://localhost:27017,localhost:27018,localhost:27019/?replicaSet=gktcsrs"
```

Check Replica Set Status

```
rs.status();
```

Check Primary/Secondary Node

```
rs.isMaster();
```

or in newer versions:

```
db.hello();
```

3. Connecting to MongoDB Atlas

```
mongosh
"mongodb+srv://user:password@cluster0.abcd.mongodb.net/myDB"
```

C. Shell Queries & Scripting Basics (15 minutes)

mongosh supports **JavaScript**, making it powerful for automation.

1. Running JavaScript Variables

```
let username = "Rahul";
db.users.find({ name: username });
```

2. Loops for Data Generation

```
for (let i = 1; i <= 10; i++) {
  db.logs.insertOne({ logId: i, ts: new Date() });
}
```

3. Writing Simple Functions

```
function findByCity(city) {
  return db.users.find({ city: city }).pretty();
}

findByCity("Pune");
```

4. Running External JavaScript File

Useful for DBA automation tasks.

Create a file: script.js

```
db.customers.insertOne({ name: "CSC Global", location:
"Bangalore" });
```

Run it:

```
mongosh script.js
```

5. Inspecting Server & DB Metadata

List DB Stats:

```
db.stats();
```

Collection Stats:

```
db.users.stats();
```

Server Build Info:

```
db.version();  
  
db.serverStatus();
```

D. LAB – Connect and Explore Collections (10–15 minutes)

This lab helps learners gain comfort with mongosh, connecting to different deployments, and exploring real data.

LAB STEPS

Step 1 – Start Shell & Connect to Local MongoDB

```
mongosh
```

Step 2 – Create Training Database

```
use csc_training
```

Step 3 – Insert Sample Data

```
db.entities.insertMany([  
  { entityName: "CSC India Pvt Ltd", city: "Bangalore",  
  active: true },  
  { entityName: "CSC Europe BV", city: "Amsterdam", active:  
  true }  
]);
```

Step 4 – Show Collections

```
show collections
```

Step 5 – Query Data

```
db.entities.find().pretty();
```

Step 6 – Connect to a Replica Set (Instructor will provide credentials)

```
mongosh  
"mongodb://localhost:27017,localhost:27018,localhost:27019/?re  
plicaSet=gktcsrs"
```

Step 7 – Explore Replica Set State

```
rs.status();
```

Step 8 – Run Queries on Replica Set

```
db.entities.find({ city: "Bangalore" });
```

Step 9 – Define a Simple Script

Create:

```
for (let i = 1; i <= 5; i++) {  
  db.audit.insertOne({ action: "TestInsert", seq: i });  
}
```

Run:

```
mongosh insert_test.js
```

🎯 LAB Outcomes

Students will be able to:

- Use mongosh confidently
- Connect to standalone & replica sets
- Run CRUD operations in shell
- Execute JS loops & scripts
- Inspect cluster health

5. Connecting to a MongoDB Database (30 min)

- Connection string concepts
- Users, roles, authentication
- Connect using mongosh, Compass, and drivers

✓ Connecting via **mongosh**, **MongoDB Compass**, and **drivers**

A. Connection String Concepts (10 minutes)

A **MongoDB connection string** (URI) tells the MongoDB client how to connect to a cluster.

1. Basic Format

mongodb://host:port

Example:

mongodb://localhost:27017

2. With Username + Password

```
mongodb://username:password@host:port/database
```

Example:

```
mongodb://adminUser:AdminPass123@localhost:27017/admin
```

3. Replica Set Connection String

Use multiple hosts + replicaSet parameter.

```
mongodb://host1:27017,host2:27018,host3:27019/?replicaSet=gktsrs
```

4. MongoDB Atlas URI (SRV format)

```
mongodb+srv://user:password@cluster.abcd.mongodb.net/mydb
```

SRV records simplify:

- Load balancing
 - Node discovery
 - SSL configuration
-

5. Common Connection String Parameters

Parameter	Purpose
replicaSet	Name of replica set
authSource	DB where credentials are stored
retryWrites	Auto retry write operations
tls=true	Enable encryption
connectTimeoutMS	Timeout for initial connection

Example with parameters:

```
mongodb://user:pass@localhost:27017/mydb?authSource=admin&retryWrites=true&tls=false
```

✓ CSC Security Requirements (Best Practices)

- Use **strong passwords & SCRAM-SHA-256**
- Use **TLS/SSL** for production
- Use **internal DNS hostnames** (not IPs)

- Always specify **authSource=admin** for superuser accounts
 - Use **application-specific users** (not admin user)
-

B. Users, Roles, and Authentication (10 minutes)

MongoDB uses **Role-Based Access Control (RBAC)**.

Authentication mechanisms

- SCRAM-SHA-1
 - SCRAM-SHA-256 (recommended)
 - LDAP (enterprise)
 - Kerberos
-

1. Create an Admin User

Run in the **admin** database:

```
use admin;

db.createUser({
  user: "cscAdmin",
  pwd: "StrongPass#2025",
  roles: [ "root" ]
});
```

2. Create Application User

Used for applications or microservices.

```
use ecommerce;

db.createUser({
  user: "appUser",
  pwd: "AppPass123",
  roles: [
    { role: "readWrite", db: "ecommerce" }
  ]
});
```

3. Common Built-In Roles

Role	Permissions
read	Read-only access
readWrite	Read + Write

Role	Permissions
dbAdmin	Create indexes, view stats
clusterAdmin	Replica set & sharding ops
root	Full access

4. Check Current User & Roles

```
db.runCommand({ connectionStatus: 1 })
```

✓ CSC Guideline for User Access**

- Admin accounts → Only DBAs
 - readWrite → Apps & services
 - read-only → Audit, reporting, compliance teams
 - clusterAdmin → Infra team
-

C. Connecting Using mongosh, Compass, and Drivers (10 minutes)

MongoDB can be accessed via:

- **mongosh (CLI)**
- **Compass (GUI)**
- **Application drivers** (Node.js, Python, Java, Go...)

1. Connect Using mongosh

✓ Connect to Local DB

```
mongosh "mongodb://localhost:27017"
```

✓ Connect with Authentication

```
mongosh
"mongodb://cscAdmin:StrongPass#2025@localhost:27017/admin"
```

✓ Connect to Replica Set

```
mongosh
"mongodb://node1:27017,node2:27018,node3:27019/?replicaSet=csc
RS"
```

✓ Connect to Atlas

```
mongosh
"mongodb+srv://cscAdmin:StrongPass#2025@cluster0.mongodb.net/m
ydb"
```

2. Connect Using MongoDB Compass (GUI)

Steps:

1. Open Compass
 2. Click “**New Connection**”
 3. Paste connection string
 4. Click “**Connect**”
-

✓ Example (Standalone)

mongodb://localhost:27017

✓ Example (Auth)

mongodb://appUser:AppPass123@localhost:27017/ecommerce?authSource=ecommerce

✓ Atlas Example

mongodb+srv://cscAdmin:<password>@cluster0.mongodb.net/

Compass automatically:

- Detects TLS
 - Displays cluster topology
 - Shows query performance
-

3. Connect using Drivers (Node.js, Python, Java)

✓ Node.js Example

```
const { MongoClient } = require("mongodb");
const uri =
"mongodb://appUser:AppPass123@localhost:27017/ecommerce";

const client = new MongoClient(uri);

async function run() {
  await client.connect();
  console.log("Connected to MongoDB");
}
run();
```

✓ Python Example

```
from pymongo import MongoClient
```

```
client =
MongoClient("mongodb://appUser:AppPass123@localhost:27017/ecommerce")
db = client.ecommerce
print(db.list_collection_names())
```

✓ Java Example

```
MongoClient client =
MongoClients.create("mongodb://appUser:AppPass123@localhost:27017/ecommerce");
MongoDatabase db = client.getDatabase("ecommerce");
```

CSC Guidance for Applications

- Use app-specific users, not admin
 - Store credentials in Vault / AWS Secret Manager
 - Always enable retryWrites
 - Use TLS=true in production
-

6. MongoDB CRUD Operations – Insert & Find (1 hr)

- insertOne(), insertMany()
- find(), findOne()
- Query filters, projections
- Lab: Perform CRUD queries on a sample dataset

CRUD = **C**reate, **R**ead, **U**update, **D**elete

This module focuses on the **Create (Insert)** and **Read (Find)** operations.

A. insertOne() and insertMany() (20 minutes)

1. insertOne()

Adds a **single document** to a collection.

Syntax:

```
db.collection.insertOne(document)
```

Example 1 – Insert a user document

```
db.users.insertOne({
```

```
    name: "Ravi Kumar",
    email: "ravi.kumar@example.com",
    city: "Bangalore"
});
```

Example 2 – Insert a CSC-style Legal Entity

```
db.entities.insertOne({
  entityName: "CSC India Pvt Ltd",
  jurisdiction: "Karnataka",
  active: true
});
```

2. insertMany()

Adds **multiple documents** at once.

Syntax:

```
db.collection.insertMany([doc1, doc2, doc3])
```

Example – Insert multiple products

```
db.products.insertMany([
  { name: "Laptop", category: "Electronics", price: 75000 },
  { name: "Mouse", category: "Electronics", price: 1200 },
  { name: "Notebook", category: "Stationery", price: 50 }
]);
```

Auto-generated `_id`

If not specified, MongoDB generates `_id: ObjectId()`

B. `find()` and `findOne()` (20 minutes)

Reading data is the most frequent operation in MongoDB.

1. `find()`

Returns a **cursor** containing all matching documents.

Syntax:

```
db.collection.find(query, projection)
```

Example – Find all documents

```
db.products.find();
```

Example – Find documents with condition

```
db.products.find({ category: "Electronics" });
```

2. `findOne()`

Returns **first matching document**.

Syntax:

```
db.collection.findOne(query)
```

Example

```
db.users.findOne({ city: "Bangalore" });
```

C. Query Filters & Projection (15 minutes)

MongoDB queries use **JSON-based filters**.

1. Equality Filter

```
db.products.find({ price: 1200 });
```

2. Comparison Operators

Operator	Meaning
\$gt	Greater than
\$gte	Greater than or equal
\$lt	Less than
\$lte	Less than or equal
\$ne	Not equal

Examples

Find products more expensive than ₹10,000:

```
db.products.find({ price: { $gt: 10000 } });
```

Find products NOT in Electronics:

```
db.products.find({ category: { $ne: "Electronics" } });
```

3. Logical Operators

Operator	Meaning
\$and	AND condition
\$or	OR condition
\$in	Matches any value in list
\$nin	Not in list

Examples

```
db.products.find({  
  $or: [  
    { category: "Electronics" },  
    { price: { $lt: 100 } }  
  ]  
});
```

4. Projection

Used to **select specific fields**.

Syntax:

```
db.collection.find(query, { field: 1, _id: 0 })
```

Examples

Return only name and price:

```
db.products.find({}, { name: 1, price: 1, _id: 0 });
```

Hide large embedded fields

```
db.entities.find({}, { officers: 0 });
```

5. Querying Embedded Fields

```
db.entities.find({ "officers.role": "Director" });
```

D. LAB – Perform CRUD Queries on a Sample Dataset (15 minutes)

Students will create a small dataset and perform CRUD reads using filters, projections, and operators.

Step 1 – Select Database

```
use csc_crud_training
```

Step 2 – Insert Sample Dataset

Insert multiple employees

```
db.employees.insertMany([
  { name: "Anita", dept: "HR", city: "Bangalore", salary: 60000 },
  { name: "Rohit", dept: "IT", city: "Pune", salary: 85000 },
  { name: "Samuel", dept: "IT", city: "Hyderabad", salary: 92000 },
  { name: "Priya", dept: "Finance", city: "Bangalore", salary: 75000 }
]);
```

Step 3 – Simple Reads

Fetch all employees:

```
db.employees.find();
```

Find employees in Bangalore:

```
db.employees.find({ city: "Bangalore" });
```

Step 4 – Comparison Filters

```
db.employees.find({ salary: { $gt: 80000 } });
```

Step 5 – Logical Filters

```
db.employees.find({
  $and: [
    { dept: "IT" },
    { salary: { $gt: 85000 } }
  ]
});
```

Step 6 – Projection Example

Show only name and salary:

```
db.employees.find({}, { name: 1, salary: 1, _id: 0 });
```

Step 7 – Embedded Document Example (Insert & Query)

```
db.customers.insertOne({  
  name: "Rakesh",  
  address: { city: "Delhi", pincode: 110001 },  
  orders: [  
    { id: "ORD001", amount: 1200 },  
    { id: "ORD002", amount: 2200 }  
  ]  
});
```

Query embedded array:

```
db.customers.find({ "orders.id": "ORD002" });
```

🎯 LAB Outcomes

Learners will understand:

- How to insert documents
 - How to use find() and findOne()
 - Filter documents using operators
 - Use projection to optimize reads
 - Query embedded and nested structures
-

7. MongoDB CRUD Operations – Replace & Delete (30 min)

- replaceOne()
- deleteOne(), deleteMany()
- Lab: Data cleanup operations

A. replaceOne() (10 minutes)

replaceOne() replaces the **entire document** except for `_id`.

✓ Syntax:

```
db.collection.replaceOne(filter, replacementDocument)
```

✓ When to use:

- Replace outdated/incorrect documents
- Overwrite entire records during migration
- Reset document structure

! Important:

- `_id` **cannot** be changed
 - If fields are missing in the new document → they are removed
-

✓ **Example 1 — Replace a user document**

Original document:

```
{  
  _id: ObjectId("..."),  
  name: "Ravi",  
  city: "Pune",  
  email: "ravi@test.com"  
}
```

Replace:

```
db.users.replaceOne(  
  { name: "Ravi" },  
  { name: "Ravi Kumar", city: "Bangalore", active: true }  
) ;
```

The field `email` is now removed (because not included in replacement).

✓ **Example 2 — CSC Legal Entity update**

```
db.entities.replaceOne(  
  { entityName: "CSC India Pvt Ltd" },  
  {  
    entityName: "CSC India Pvt Ltd",  
    jurisdiction: "Karnataka",  
    complianceStatus: "Active",  
    updatedBy: "adminUser"  
  }  
) ;
```

B. `deleteOne()` and `deleteMany()` (10 minutes)

✓ **1. `deleteOne()`**

Deletes **the first** document matching the filter.

```
db.collection.deleteOne({ field: value })
```

Example:

```
db.products.deleteOne({ name: "Notebook" });
```

✓ 2. deleteMany()

Deletes **all documents** matching the filter.

```
db.collection.deleteMany({ category: "Expired" });
```

✓ Example: Delete inactive CSC compliance logs

```
db.complianceLogs.deleteMany({ status: "obsolete" });
```

✓ Example: Delete employees from a specific city

```
db.employees.deleteMany({ city: "Delhi" });
```

! Safety Tips for Delete Operations

- Always run find() with same filter before delete
 - Ensure backups exist
 - For high-risk operations, use transactions (in replica sets/sharded clusters)
 - Avoid empty filter {} unless intentional
-

C. LAB – Data Cleanup Operations (10 minutes)

Students will create sample data, perform replaceOne(), and cleanup using deletes.

Step 1 – Create Collection

```
use csc_crud_lab;

db.inventory.insertMany([
  { item: "Laptop", qty: 10, category: "Electronics" },
  { item: "Mouse", qty: 50, category: "Electronics" },
  { item: "Notebook", qty: 100, category: "Stationery" },
```

```
{ item: "OldFile", qty: 0, category: "Obsolete" }  
] );
```

Step 2 – Replace Document

Replace “Laptop” with new structure:

```
db.inventory.replaceOne(  
  { item: "Laptop" },  
  { item: "Laptop", qty: 8, category: "Electronics", updated:  
true }  
) ;
```

Step 3 – Delete One Obsolete Item

```
db.inventory.deleteOne({ item: "OldFile" }) ;
```

Step 4 – Delete Many Stationery Items

```
db.inventory.deleteMany({ category: "Stationery" }) ;
```

Step 5 – Display Final Collection

```
db.inventory.find().pretty();
```

Expected result:

- Laptop remains (replaced version)
 - Mouse remains
 - Notebook deleted
 - OldFile deleted
-

🎯 LAB Outcomes:

Learners will understand:

- How replaceOne() overwrites documents
 - Safe delete operations
 - Cleanup patterns for real-world DB maintenance
 - Differences between deleteOne() and deleteMany()
-

8. Modifying Query Results (1 hr)

- `sort()`, `skip()`, `limit()`
 - Aggregation preview
 - Query optimization basics
 - Lab: Apply query modifiers for performance
-

MODULE 8: Modifying Query Results (1 hr)

MongoDB allows result modification using:

- `sort()`
- `limit()`
- `skip()`

These operations help optimize read queries, especially for pagination and reporting.

A. `sort()`, `limit()`, `skip()`

These functions modify **how results are returned**, not how they are stored.

1. `sort()`

Sorts documents based on one or more fields.

✓ Syntax:

```
db.collection.find().sort({ field: 1 }) // ascending
```

```
db.collection.find().sort({ field: -1 }) // descending
```

✓ Example – Sort employees by salary (highest first)

```
db.employees.find().sort({ salary: -1 });
```

✓ Example – Multi-field sort

Sort by department ascending, then salary descending:

```
db.employees.find().sort({ dept: 1, salary: -1 });
```

⚠️ Performance Note:

- Sorting **without an index** can cause blocking operations.
 - Create an index on sorted fields for faster results.
-

2. limit()

Restricts number of documents returned.

✓ Syntax:

```
db.collection.find().limit(5)
```

Example – Get top 5 highest salaries:

```
db.employees.find().sort({ salary: -1 }).limit(5);
```

3. skip()

Skip first N documents.

✓ Syntax:

```
db.collection.find().skip(10)
```

Example – Pagination:

```
db.employees.find().sort({ name: 1 }).skip(20).limit(10);
```

⭐ Pagination Pattern (skip & limit)

Used for:

- CSC dashboards
- Paginated corporate records
- Trademark search results

Typical pattern:

```
db.collection.find()  
  .sort({ createdAt: -1 })  
  .skip(page * pageSize)  
  .limit(pageSize)
```

B. Aggregation Preview (10 minutes)

Aggregation pipeline = MongoDB's equivalent of SQL GROUP BY + WHERE + ORDER + FUNCTIONS.

★ Common Stages:

✓ 1. \$match – Filter documents

```
{ $match: { dept: "IT" } }
```

✓ 2. \$group – Group and aggregate

```
{ $group: { _id: "$dept", totalSalary: { $sum: "$salary" } } }
```

✓ 3. \$sort – Sort aggregated output

```
{ $sort: { totalSalary: -1 } }
```

★ Example – Department salary totals

```
db.employees.aggregate([
  { $match: {} },
  { $group: { _id: "$dept", total: { $sum: "$salary" } } },
  { $sort: { total: -1 } }
]);
```

★ Why Aggregation Matters for CSC?

- Generating compliance dashboards
 - Risk scoring
 - Trademark classification reports
 - Officer distribution summaries
-

C. Query Optimization Basics (10 minutes)

Efficient queries reduce latency and resource usage.

★ 1. Use Indexes on Sorted or Filtered Fields

Sorting without an index forces **in-memory sort** → slow.

Example:

```
db.employees.find().sort({ salary: -1 });
```

Needs index:

```
db.employees.createIndex({ salary: -1 });
```

★ 2. Avoid skip() for large offsets

skip(1000000) becomes slow.

Alternatives:

- Range queries
 - Bookmark-based pagination
-

★ 3. Use Projection to Reduce Data Size

Return only required fields:

```
db.employees.find({}, { name: 1, dept: 1, _id: 0 })
```

★ 4. Analyze Query Plan with explain()

```
db.employees.find({ dept: "IT" }).explain("executionStats");
```

Key terms:

- **COLLSCAN** → full collection scan (slow)
 - **IXSCAN** → index scan (fast)
 - **nReturned**
 - **executionTimeMillis**
-

★ 5. Avoid OR conditions when possible

Prefer:

```
$in: ["Pune", "Bangalore"]
```

instead of:

```
$or: [{ city: "Pune" }, { city: "Bangalore" }]
```

D. LAB – Apply Query Modifiers for Performance (15 minutes)

Goal:

Use sort(), limit(), skip() and analyze performance with explain().

Step 1 – Dataset Setup

```
use csc_querylab;

for (let i = 1; i <= 50000; i++) {
  db.employees.insertOne({
    empId: i,
    name: "Employee" + i,
    dept: ["IT", "HR", "Finance", "Legal"][i % 4],
    salary: Math.floor(Math.random() * 90000) + 30000,
    city: ["Bangalore", "Pune", "Chennai", "Delhi"][i % 4]
  });
}
```

Step 2 – Query Without Index

```
db.employees.find({ city: "Bangalore" }).sort({ salary: -1 })
  .limit(5);
```

Step 3 – Check Query Plan

```
db.employees.find({ city: "Bangalore" })
  .sort({ salary: -1 })
  .limit(5)
  .explain("executionStats");
```

Look for:

- COLLSCAN
 - High nScannedDocs
 - High executionTimeMillis
-

Step 4 – Create Index

```
db.employees.createIndex({ city: 1, salary: -1 });
```

Step 5 – Run Query Again

```
db.employees.find({ city: "Bangalore" })
  .sort({ salary: -1 })
  .limit(5)
  .explain("executionStats");
```

Expected:

- IXSCAN
 - Fewer scanned docs
 - Lower execution time
-

Step 6 – Pagination Exercise

```
db.employees.find({ dept: "IT" })
  .sort({ empId: 1 })
  .skip(100)
  .limit(10);
```

🎯 Lab Outcomes:

Learners will understand:

- How sorting affects performance
 - How skip & limit create pagination
 - How to use explain()
 - Indexes needed for optimal query performance
 - Real-world query tuning strategies
-

Day1 Assignment

- CRUD tasks on sample collections
 - Short quiz (MCQ + practical)
-

A — CRUD TASKS ON SAMPLE COLLECTIONS

Use a new database:

```
use day1_assignment;
```

1 Create Collections & Insert Documents

Task 1. Insert sample employees

Insert **5** employee documents with fields:

- name
- department
- city
- salary
- active (boolean)

Use **insertMany()**.

Task 2. Insert CSC-style legal entities

Insert **3** documents with fields:

- entityName
- jurisdiction
- status
- officers (array of embedded documents)
- lastFiled (Date)

Use **insertOne()**.

2 Find & Filter Operations

Task 3. Read operations

Write queries using **find()** to:

- a) Get all employees from Bangalore
 - b) Get employees with salary > 75,000
 - c) Get legal entities with status = "Active"
 - d) Find officers whose role = "Director" using embedded query
 - e) Project only name and salary fields for employees
-

3 Update & Replace Operations

Task 4. Update employee city

Update city of **one employee** from Bangalore → Pune.

Use:

```
updateOne({ ... }, { $set: { ... } })
```

Task 5. Replace an entire entity document

Replace one entity document with a **new structure** (except `_id`).

Use:

```
replaceOne()
```

4 Delete Operations

Task 6. Delete inactive employees

Delete employees where:

```
{ active: false }
```

Task 7. Delete officers with specific role inside an entity

(Use `$pull` operator)

Example:

```
$pull: { officers: { role: "Secretary" } }
```

PART B — QUERY MODIFIER TASKS

(sort, skip, limit)

Task 8. Sort employees by salary, descending

```
sort({ salary: -1 })
```

Task 9. Return the top 3 highest-paid employees

Use sort + limit.

Task 10. Apply pagination

Fetch **page 2** of results, page size = 2 employees.

Use:

```
sort().skip().limit()
```

PRACTICAL QUESTIONS (Short Coding Problems) (20 minutes)

Q1. Insert 100 sample audit logs using a loop.

Each log should contain:

- logId
 - user
 - ts (timestamp)
-

Q2. Query audit logs created today.

Q3. Write a query to fetch employees with salary BETWEEN 60,000 and 90,000.

Use \$gte and \$lte.

Q4. Write a query to return only employees from IT or HR department.

Use \$in.

Q5. Write a projection query to hide _id and show only name + department.

Q6. Find all customers whose order total > ₹10,000 using embedded query.

Dataset example:

```
{  
  name: "...",  
  orders: [  
    { id: "O1", amount: 5000 },  
    { id: "O2", amount: 6000 }  
  ]  
}
```

Q7. Explain the difference between find() and findOne().

Q8. Why is limit() important for optimization? Give 1 example.

Day1 MongoDB Admin SHORT MCQ QUIZ (10 Questions) (10 minutes)

1. Which command inserts multiple documents?

- a) insertAll()
 - b) insertMany()
 - c) createMany()
 - d) saveMany()
-

2. Which query returns only the first matching document?

- a) find()
- b) findFirst()

c) `findOne()`

d) `findTop()`

3. Which operator finds values greater than 100?

a) `$gt`

b) `$gte`

c) `$greater`

d) `>`

4. What does projection control?

a) Which documents to insert

b) Which fields to show

c) Which documents to delete

d) Which database to create

5. `replaceOne()` removes missing fields in the new document.

a) True

b) False

6. `deleteMany()` deletes:

a) First matching document

b) All matching documents

c) Documents with `_id` only

d) None

7. `sort({salary: -1})` means:

- a) Sort by salary ascending
 - b) Sort by salary descending
 - c) Group by salary
 - d) Remove salary field
-

8. `skip(5)` means:

- a) Delete 5 documents
 - b) Skip the first 5 documents
 - c) Sort 5 documents
 - d) Insert 5 documents
-

9. Which operator matches ANY of the listed values?

- a) `$match`
 - b) `$in`
 - c) `$or`
 - d) `$contains`
-

10. `explain("executionStats")` helps understand:

- a) Document size
 - b) Query plan + performance
 - c) Password encryption
 - d) Replica set members
-