

Introduction to Replication in MongoDB

- Replica set architecture
- Primary, Secondary, Arbiter
- Elections & failover
- Read preference & write concerns
- Lab: Set up a 3-node replica set locally
- Lab: Test failover + observe elections

1. Set up a **3-node replica set locally**
 2. **Test failover** and observe elections, read preferences, and write concerns
-

0. CSC Story / Use Case (use this narrative in training)

CSC runs a **Compliance & Filings Platform** for thousands of corporate clients.

They **cannot afford downtime** during US year-end filing season.

- Production runs on a **3-node MongoDB replica set** (Linux or Atlas).
- Bangalore SRE team is responsible for **high availability** and **disaster recovery**.
- Mumbai team builds reporting and analytics that must read from **secondaries** without impacting primary.

To do this safely, everyone must understand:

- **Replica set architecture**
 - **Primary / Secondary / Arbiter**
 - **Elections & failover**
 - **Read preferences & write concerns**
 - How to **set up and break** a replica set in the lab, *before* doing anything in production.
-

1. Concepts – Replica Set Basics (Short & Clear)

1.1 Replica Set Architecture

A **replica set** is a group of mongod processes that maintain the **same data set**.

Typical CSC production pattern:

- **Primary** – handles all writes and default reads.
- **Secondaries (1–2 or more)** – replicate data from primary; can serve **read-only** traffic.
- (Optional) **Arbiter** – only votes in elections; **no data**.

Architecture example (3-node):

- csc-rs-1 – PRIMARY
 - csc-rs-2 – SECONDARY
 - csc-rs-3 – SECONDARY or ARBITER (depending on design)
-

1.2 Primary, Secondary, Arbiter

- **Primary**
 - Accepts all writes.
 - Replicates changes to secondaries via **oplog**.
- **Secondary**
 - Copies data from primary (oplog).
 - Can be used for **reads**, backups, reporting.
- **Arbiter**
 - No data.
 - Only participates in **elections** to avoid ties (odd number of votes).

CSC-style use:

- Primary: in Bangalore or primary DC
 - Secondaries: one more in same DC, another in DR region or Mumbai DC
 - Arbiter: sometimes used in small clusters where you want 3 votes but only 2 data-bearing nodes.
-

1.3 Elections & Failover

If a **primary goes down**:

1. Replica set members detect primary is unavailable.
2. Eligible secondary calls an **election**.
3. New primary is elected based on:
 - Priority
 - Data “freshness”

- Voting majority

Failover is **automatic** (within a few seconds) → this is what keeps CSC app **online**.

1.4 Read Preferences & Write Concerns

- **Read Preference**
 - primary (default) – safest, latest data.
 - secondary – offload reads (e.g., reporting in Mumbai).
 - primaryPreferred, secondaryPreferred, nearest – hybrid strategies.
- **Write Concern**
 - w: 1 – acknowledge from primary only.
 - w: "majority" – acknowledge when majority of nodes have written (safer).
 - j: true – wait for journal, improving durability.

CSC production typical choice:

- Most critical writes: writeConcern: { w: "majority", j: true }
 - Some internal logs/metrics: w: 1
-

2. Lab 1 – Set up a 3-node Replica Set LOCALLY

We'll simulate a 3-node replica set on a **single machine** using **three different data folders and ports**.

We'll show the setup generically, then highlight OS-specific differences:

- **Linux / Mac** – similar shell commands.
- **Windows** – use cmd or PowerShell.

For simplicity, we'll name:

- Replica set name: cscReplSet
- Ports: 27017, 27018, 27019
- Data dirs: data/rs1, data/rs2, data/rs3

Important: Make sure mongod is installed and on PATH.

2.1 Prepare Data Directories

Linux / Mac

```
mkdir -p ~/csc-repl/data/rs1 ~/csc-repl/data/rs2 ~/csc-repl/data/rs3
```

Windows (PowerShell)

```
mkdir C:\csc-repl\data\rs1  
mkdir C:\csc-repl\data\rs2  
mkdir C:\csc-repl\data\rs3
```

2.2 Start 3 mongod Instances

Linux / Mac

Open **three terminals**, one for each node:

Node 1:

```
mongod --replSet cscReplSet --port 27017 --dbpath ~/csc-repl/data/rs1 --  
bind_ip 127.0.0.1
```

Node 2:

```
mongod --replSet cscReplSet --port 27018 --dbpath ~/csc-repl/data/rs2 --  
bind_ip 127.0.0.1
```

Node 3:

```
mongod --replSet cscReplSet --port 27019 --dbpath ~/csc-repl/data/rs3 --  
bind_ip 127.0.0.1
```

Windows (cmd or PowerShell)

Again, three separate windows:

```
mongod --replSet cscReplSet --port 27017 --dbpath C:\csc-repl\data\rs1 --  
bind_ip 127.0.0.1  
mongod --replSet cscReplSet --port 27018 --dbpath C:\csc-repl\data\rs2 --  
bind_ip 127.0.0.1
```

```
mongod --replSet cscReplSet --port 27019 --dbpath C:\csc-repl\data\rs3 --  
bind_ip 127.0.0.1
```

At this point, 3 standalone mongods are running with the **same replica set name**.

2.3 Initialize the Replica Set

Connect to **node 1** using mongosh:

```
mongosh "mongodb://127.0.0.1:27017"
```

Run:

```
rs.initiate({  
  _id: "cscReplSet",  
  members: [  
    { _id: 0, host: "127.0.0.1:27017" },  
    { _id: 1, host: "127.0.0.1:27018" },  
    { _id: 2, host: "127.0.0.1:27019" }  
  ]  
})
```

Check status:

```
rs.status()
```

You should see:

- One node as **PRIMARY**
- Two nodes as **SECONDARY**

CSC use case explanation

This mimics CSC's **3-node production replica set**.

In reality, instead of ports on one machine, these hosts would be:

- mongo1.bangalore.csc.com:27017
 - mongo2.bangalore.csc.com:27017
 - mongo3.mumbai.csc.com:27017
-

2.4 Insert & Read Data from Replica Set

Still on primary:

```
use csc_compliance

db.filings.insertOne({
  entity_id: 1001,
  filing_type: "Annual Report",
  state: "DE",
  status: "OPEN",
  amount: 5000
})
```

Check:

```
db.filings.find().pretty()
```

Now connect to a **secondary** to see replication (for test, we'll enable secondary reads):

```
mongosh "mongodb://127.0.0.1:27018"
```

By default, secondaries don't allow reads, so run:

```
rs.slaveOk() // for test only (legacy name; still works)
use csc_compliance
db.filings.find().pretty()
```

You should see the same document → replication working.

CSC use case explanation

Mumbai reporting service reads from **secondaries** to generate heavy monthly compliance reports without overloading the primary.

3. Lab 2 – Elections, Failover, Read Preference, Write Concerns

Now we'll **break things on purpose** and see how the replica set behaves.

3.1 Trigger an Election (Primary Step Down)

Connect to the **PRIMARY** (usually 27017):

```
mongosh "mongodb://127.0.0.1:27017"
```

Run:

```
rs.status() // confirm this is PRIMARY  
rs.stepDown(60) // step down for 60 seconds
```

What happens:

- This node becomes **SECONDARY**.
- Another node (27018 or 27019) becomes **PRIMARY**.

Check from a different node:

```
mongosh "mongodb://127.0.0.1:27018"  
rs.isMaster()
```

(Or in newer versions, db.hello().)

You should see:

```
{ isWritablePrimary: true, ... }
```

CSC use case explanation

If the primary server in Bangalore goes down, the secondary in another Bangalore rack or Mumbai DC can **take over automatically**, so client applications keep working.

3.2 Simulate Primary Crash

Instead of stepDown, you can **kill the mongod process** of the current primary window (Ctrl+C).

Watch rs.status() from another node and see election logs.

3.3 Read Preference – Reading from Secondary

Let's connect using a **replica set connection string** and set read preference.

```
mongosh  
"mongodb://127.0.0.1:27017,127.0.0.1:27018,127.0.0.1:27019/csc_compliance?replicaSet=cscReplSet"
```

By default:

```
db.getMongo().getReadPref() // primary
```

Change read preference to secondary (for test):

```
db.getMongo().setReadPref("secondary")  
db.filings.find().limit(1)
```

Mongo will route reads to a secondary.

✓ CSC use case explanation

Reporting jobs from Mumbai can be configured with readPreference: secondary to **offload read traffic** from the primary.

3.4 Write Concerns – Testing Safety Levels

In mongosh:

```
db.filings.insertOne(  
  { entity_id: 1002, filing_type: "Franchise Tax", state: "CA", status: "OPEN", amount: 8000 },  
  { writeConcern: { w: 1 } })
```

This only waits for **primary**.

Now try majority:

```
db.filings.insertOne(  
  { entity_id: 1003, filing_type: "Annual Report", state: "NY", status:  
  "OPEN", amount: 12000 },  
  { writeConcern: { w: "majority", j: true } }  
)
```

This waits until a **majority of nodes** have the write and it's **journaled**.

CSC use case explanation

For critical filings, CSC uses w: "majority", j: true so a write is guaranteed to be on at least 2 data-bearing nodes *and* durable on disk before the application proceeds.

4. Replication in MongoDB Atlas (Concept + Quick Hands-On)

In Atlas, **replication is built-in**. A standard Atlas cluster is already a replica set.

4.1 View Replica Set Members

1. Log in to **Atlas**.
2. Go to your project → open cluster (e.g., csc-compliance-prod).
3. Click **Connect** → **Connect with MongoDB Shell** to get the mongosh URI.
4. Connect:

```
mongosh  
"mongodb+srv://csc_app_user@cluster0.xxxxxx.mongodb.net/csc_compliance"
```

5. Check replica set config:

```
rs.status()
```

You'll see the Atlas nodes (3+ members).

4.2 Test Read Preference on Atlas

In mongosh (Atlas):

```
use csc_compliance
db.getMongo().getReadPref() // should show primary

db.getMongo().setReadPref("secondaryPreferred")
db.filings.find().limit(1)
```

In real apps (Java, Node, .NET), you set read preference in the **driver connection string** or code.

4.3 Test Failover on Atlas (Carefully – Use Non-Prod)

In Atlas UI:

1. Go to your cluster.
2. Click the ... menu → **Test Failover** (or similar action, depending on UI and permissions).
3. Atlas will **step down** the primary and elect a new one.
4. In mongosh, watch:

```
while (true) {
    printjson(rs.isMaster());
    sleep(2000);
}
```

You'll see isWritablePrimary flip from one node to another.

CSC use case explanation

This is how CSC tests **application resilience**: the app should automatically reconnect to the new primary during Atlas failover tests.

5. How It Looks Across Platforms (Quick Summary Table)

Platform	What You Do for This Lab
Linux	Run 3 mongod processes with different ports and data dirs, initialize replica set, test elections.
Windows	Same idea, but with C:\csc-repl paths and cmd/PowerShell windows.
Mac	Same as Linux, usually installed via Homebrew, using ~/csc-repl.
Atlas	Replica set is auto-managed. You inspect with rs.status(), set readPreference in the driver, and use “Test Failover” in UI.

6. Suggested Exercise Flow for a 90-Minute Session

1. **(15 min)** Explain replica set architecture + primary/secondary/arbitrer.
 2. **(20–30 min)** Lab 1: Set up 3-node replica set (Windows/Mac/Linux).
 3. **(20 min)** Lab 2: Insert data, test secondary reads, step down primary, observe elections.
 4. **(10–15 min)** Explain read preferences & write concerns; try a couple of combinations.
 5. **(15–20 min)** Atlas: Show rs.status(), read preferences, and a test failover using Atlas UI (non-prod cluster).
-